

21051923 邵屹杰

(代码环境在文末说明)

tf2.0-exercise 代码与结果

实现softmax函数

[illegible]

实现sigmoid函数

[illegible]

实现 softmax 交叉熵loss函数

```
1 def softmax_ce(x, label):
2     #####
3     '''实现 softmax 交叉熵loss函数, 不允许用tf自带的softmax_cross_entropy函数'''
4     #####
5     log_softmax_x = np.log(x)
6     loss = tf.reduce_mean(-np.sum(label * log_softmax_x, axis=-1))
7     return loss
8
9 test_data = np.random.normal(size=[10, 5])
10 prob = tf.nn.softmax(test_data)
11 label = np.zeros_like(test_data)
12 label[np.arange(10), np.random.randint(0, 5, size=10)]=1.
13
14 ((tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(label, test_data))
15  - softmax_ce(prob, label))*2 < 0.0001).numpy()
```

[54]

... True

实现 sigmoid 交叉熵loss函数

```
1 def sigmoid_ce(x, label):
2     #####
3     '''实现 sigmoid 交叉熵loss函数, 不允许用tf自带的softmax_cross_entropy函数'''
4     #####
5     loss = tf.convert_to_tensor(-np.mean(label * np.log(x) + (1 - label) * np.log(1 - x)))
6     return loss
7
8 test_data = np.random.normal(size=[10])
9 prob = tf.nn.sigmoid(test_data)
10 label = np.random.randint(0, 2, 10).astype(test_data.dtype)
11 print(label)
12
13 ((tf.reduce_mean(tf.nn.sigmoid_cross_entropy_with_logits(label, test_data))- sigmoid_ce(prob, label))*2 < 0.0001).numpy()
```

[55]

... [1. 1. 0. 1. 1. 0. 1. 1. 1. 1.]

True

tutorial_minst_fnn-numpy-exercise 代码填空部分与结果

```
11
12 def backward(self, grad_y):
13     '''
14     x: shape(N, d)
15     w: shape(d, d')
16     grad_y: shape(N, d')
17     '''
18     x = self.mem['x']
19     W = self.mem['W']
20
21     #####
22     '''计算矩阵乘法的对应的梯度'''
23     #####
24     grad_W = np.dot(x.T, grad_y)
25     grad_x = np.dot(grad_y, W.T)
26     return grad_x, grad_W
27
```

```

29 class Relu:
30     def __init__(self):
31         self.mem = {}
32
33     def forward(self, x):
34         self.mem['x']=x
35         return np.where(x > 0, x, np.zeros_like(x))
36
37     def backward(self, grad_y):
38         '''
39         grad_y: same shape as x
40         '''
41         #####
42         '''计算relu 激活函数对应的梯度'''
43         #####
44         x = self.mem['x']
45         grad_x = np.where(x > 0, grad_y, np.zeros_like(grad_y))
46         return grad_x
47

```

实际训练

```

1 train_data, test_data = mnist_dataset()
2 train_label = np.zeros(shape=[train_data[0].shape[0], 10])
3 test_label = np.zeros(shape=[test_data[0].shape[0], 10])
4 train_label[np.arange(train_data[0].shape[0]), np.array(train_data[1])] = 1.
5 test_label[np.arange(test_data[0].shape[0]), np.array(test_data[1])] = 1.
6
7 for epoch in range(50):
8     loss, accuracy = train_one_step(model, train_data[0], train_label)
9     print('epoch', epoch, ' : loss', loss, ' ; accuracy', accuracy)
10 loss, accuracy = test(model, test_data[0], test_label)
11
12 print('test loss', loss, ' ; accuracy', accuracy)

```

```

... epoch 0 : loss 23.19814963363836 ; accuracy 0.11136666666666667
epoch 1 : loss 21.652869263919637 ; accuracy 0.15455
epoch 2 : loss 19.18432926774554 ; accuracy 0.2287
epoch 3 : loss 17.410781618645153 ; accuracy 0.30536666666666667
epoch 4 : loss 16.176548322132227 ; accuracy 0.33873333333333333
epoch 5 : loss 15.278108643437454 ; accuracy 0.37823333333333333
epoch 6 : loss 14.500877858453324 ; accuracy 0.40468333333333334
epoch 7 : loss 13.928311874486658 ; accuracy 0.4252
epoch 8 : loss 13.433952620260188 ; accuracy 0.44478333333333333
epoch 9 : loss 13.37007614396068 ; accuracy 0.4469
epoch 10 : loss 12.629945948951876 ; accuracy 0.48103333333333333
epoch 11 : loss 12.25486790822371 ; accuracy 0.48863333333333333
epoch 12 : loss 12.094745666993942 ; accuracy 0.5056
epoch 13 : loss 11.59343101759849 ; accuracy 0.51683333333333334
epoch 14 : loss 11.308935254491034 ; accuracy 0.53545
epoch 15 : loss 10.98401811658683 ; accuracy 0.54165
epoch 16 : loss 10.734419954092768 ; accuracy 0.5577
epoch 17 : loss 10.481756352318058 ; accuracy 0.56266666666666666
epoch 18 : loss 10.264747324050068 ; accuracy 0.57623333333333334

```

Jupyter > exercise-master > chap4_simple neural network > tutorial_minst_fnn-numpy-exercise.ipynb > M*准备数据

十 代码 十 Markdown | 全部运行 清除所有输出 重启 变量 大纲 ...

```
epoch 14 : loss 11.308935254491034 ; accuracy 0.53545
epoch 15 : loss 10.98401811658683 ; accuracy 0.54165
epoch 16 : loss 10.734419954092768 ; accuracy 0.5577
epoch 17 : loss 10.481756352318058 ; accuracy 0.5626666666666666
epoch 18 : loss 10.264747324050068 ; accuracy 0.5762333333333334
epoch 19 : loss 10.066426266778773 ; accuracy 0.5808833333333333
epoch 20 : loss 9.916854678473376 ; accuracy 0.5902333333333334
epoch 21 : loss 9.760156675082035 ; accuracy 0.5933833333333334
epoch 22 : loss 9.64539379343153 ; accuracy 0.6008833333333333
epoch 23 : loss 9.499947779628059 ; accuracy 0.6042166666666666
epoch 24 : loss 9.399045832591124 ; accuracy 0.6115666666666667
epoch 25 : loss 9.277739763641918 ; accuracy 0.6139666666666667
epoch 26 : loss 9.184297644622024 ; accuracy 0.6196166666666667
epoch 27 : loss 9.08683964163651 ; accuracy 0.6221833333333333
epoch 28 : loss 9.007161831604638 ; accuracy 0.627
epoch 29 : loss 8.92218185193055 ; accuracy 0.6291833333333333
epoch 30 : loss 8.8521656727106 ; accuracy 0.63355
epoch 31 : loss 8.767432663554414 ; accuracy 0.6352333333333333
epoch 32 : loss 8.697078195733697 ; accuracy 0.6393833333333333
epoch 33 : loss 8.60378116114664 ; accuracy 0.6408833333333334
epoch 34 : loss 8.5196229005761 ; accuracy 0.6452166666666667
epoch 35 : loss 8.390902509682515 ; accuracy 0.6475166666666666
epoch 36 : loss 8.210584842081415 ; accuracy 0.6540333333333334
epoch 37 : loss 7.8570336966631436 ; accuracy 0.6609333333333334
epoch 38 : loss 7.368278578760774 ; accuracy 0.67505
epoch 39 : loss 6.877058237281882 ; accuracy 0.6916166666666667
epoch 40 : loss 6.591557242401064 ; accuracy 0.70635
epoch 41 : loss 6.446095209271721 ; accuracy 0.71115
epoch 42 : loss 6.5189907522508586 ; accuracy 0.71125
epoch 43 : loss 6.68820329794945 ; accuracy 0.7043333333333334
epoch 44 : loss 6.438024798546668 ; accuracy 0.71595
epoch 45 : loss 6.238570298895335 ; accuracy 0.7244666666666667
epoch 46 : loss 5.958566064263271 ; accuracy 0.7365666666666667
epoch 47 : loss 5.904002112404578 ; accuracy 0.7384
epoch 48 : loss 5.761329829070316 ; accuracy 0.7455666666666667
epoch 49 : loss 5.7170502127872185 ; accuracy 0.7468333333333333
test loss 5.49741895973345 ; accuracy 0.7569
```

tutorial_minst_fnn-tf2.0-exercise 代码填空与结果

```
1 class myModel:
2     def __init__(self):
3         #####
4         '''声明模型对应的参数'''
5         #####
6         num = 1000
7         self.W1 = tf.Variable(shape=[28 * 28, num], dtype=tf.float32,
8                                initial_value=tf.random.uniform(shape=[28 * 28, num], minval=-0.1, maxval=0.1))
9         self.b1 = tf.Variable(shape=[num], dtype=tf.float32, initial_value=tf.zeros(num))
10        self.W2 = tf.Variable(shape=[num, 10], dtype=tf.float32,
11                               initial_value=tf.random.uniform(shape=[num, 10], minval=-0.1, maxval=0.1))
12        self.b2 = tf.Variable(shape=[10], dtype=tf.float32, initial_value=tf.zeros(10))
13
14
15    def __call__(self, x):
16        #####
17        '''实现模型函数体，返回未归一化的logits'''
18        #####
19        flat_x = tf.reshape(x, shape=[-1, 28 * 28])
20        h1 = tf.nn.relu(tf.matmul(flat_x, self.W1) + self.b1)
21        logits = tf.matmul(h1, self.W2) + self.b2
22        return logits
23
24 model = myModel()
25
26 optimizer = optimizers.Adam()
```

实际训练

```
1 train_data, test_data = mnist_dataset()
2 for epoch in range(50):
3     loss, accuracy = train_one_step(model, optimizer,
4                                     tf.constant(train_data[0], dtype=tf.float32),
5                                     tf.constant(train_data[1], dtype=tf.int64))
6     print('epoch', epoch, ': loss', loss.numpy(), '; accuracy', accuracy.numpy())
7 loss, accuracy = test(model,
8                       tf.constant(test_data[0], dtype=tf.float32),
9                       tf.constant(test_data[1], dtype=tf.int64))
10
11 print('test loss', loss.numpy(), '; accuracy', accuracy.numpy())
```

[63]

```
... epoch 0 : loss 2.3905141 ; accuracy 0.20203333
epoch 1 : loss 2.337947 ; accuracy 0.21081667
epoch 2 : loss 2.291935 ; accuracy 0.21888334
epoch 3 : loss 2.2509954 ; accuracy 0.22891666
epoch 4 : loss 2.214041 ; accuracy 0.23996666
epoch 5 : loss 2.1802578 ; accuracy 0.25273332
epoch 6 : loss 2.1490245 ; accuracy 0.26703334
epoch 7 : loss 2.119862 ; accuracy 0.28203332
epoch 8 : loss 2.0923965 ; accuracy 0.2974
epoch 9 : loss 2.0663376 ; accuracy 0.31305
epoch 10 : loss 2.0414567 ; accuracy 0.33028334
epoch 11 : loss 2.017575 ; accuracy 0.34721667
epoch 12 : loss 1.9945508 ; accuracy 0.36391667
epoch 13 : loss 1.9722735 ; accuracy 0.38
epoch 14 : loss 1.9506553 ; accuracy 0.39553332
epoch 15 : loss 1.9296277 ; accuracy 0.4103
epoch 16 : loss 1.9091341 ; accuracy 0.4246
epoch 17 : loss 1.8891299 ; accuracy 0.43905
epoch 18 : loss 1.8695794 ; accuracy 0.45206666
epoch 19 : loss 1.8504533 ; accuracy 0.4649
epoch 20 : loss 1.8317275 ; accuracy 0.47716665
```

```
epoch 18 : loss 1.8695794 ; accuracy 0.45206666
epoch 19 : loss 1.8504533 ; accuracy 0.4649
epoch 20 : loss 1.8317275 ; accuracy 0.47716665
epoch 21 : loss 1.8133824 ; accuracy 0.48941666
epoch 22 : loss 1.7954013 ; accuracy 0.50123334
epoch 23 : loss 1.7777705 ; accuracy 0.51241666
epoch 24 : loss 1.7604773 ; accuracy 0.52271664
epoch 25 : loss 1.7435107 ; accuracy 0.53225
epoch 26 : loss 1.7268608 ; accuracy 0.54145
epoch 27 : loss 1.7105186 ; accuracy 0.55055
epoch 28 : loss 1.6944765 ; accuracy 0.55943334
epoch 29 : loss 1.6787268 ; accuracy 0.56815
epoch 30 : loss 1.6632627 ; accuracy 0.57738334
epoch 31 : loss 1.6480784 ; accuracy 0.5857667
epoch 32 : loss 1.633168 ; accuracy 0.5934333
epoch 33 : loss 1.618525 ; accuracy 0.6006333
epoch 34 : loss 1.6041437 ; accuracy 0.6073
epoch 35 : loss 1.5900184 ; accuracy 0.61406666
epoch 36 : loss 1.5761435 ; accuracy 0.61981666
epoch 37 : loss 1.5625143 ; accuracy 0.6259
epoch 38 : loss 1.5491256 ; accuracy 0.63235
epoch 39 : loss 1.535973 ; accuracy 0.63875
epoch 40 : loss 1.5230509 ; accuracy 0.64395
epoch 41 : loss 1.5103546 ; accuracy 0.64893335
epoch 42 : loss 1.4978802 ; accuracy 0.65393335
epoch 43 : loss 1.4856224 ; accuracy 0.65885
epoch 44 : loss 1.4735763 ; accuracy 0.66326666
epoch 45 : loss 1.4617378 ; accuracy 0.66735
epoch 46 : loss 1.4501021 ; accuracy 0.6713833
epoch 47 : loss 1.4386652 ; accuracy 0.67571664
epoch 48 : loss 1.4274224 ; accuracy 0.6794
epoch 49 : loss 1.4163706 ; accuracy 0.68295
test loss 1.3832035 ; accuracy 0.6959
```

运行环境 mytensor

Package	Version

absl-py	1.4.0
astroid	2.15.0
astunparse	1.6.3
attrs	22.2.0
autopep8	1.6.0
backcall	0.2.0
cached-property	1.5.2
cachetools	5.3.0
certifi	2022.12.7
charset-normalizer	3.0.1
chex	0.1.5
colorama	0.4.6
cycler	0.11.0
debugpy	1.5.1
decorator	5.1.1
dill	0.3.6
dm-pix	0.4.0
dm-tree	0.1.8
docstring-to-markdown	0.11
entrypoints	0.4
etils	0.9.0
exceptiongroup	1.1.0
flake8	5.0.4
flatbuffers	23.3.3
flax	0.6.4
fonttools	4.38.0
gast	0.4.0
gin-config	0.5.0
google-auth	2.16.2
google-auth-oauthlib	0.4.6
google-pasta	0.2.0
grpcio	1.51.3
h5py	3.8.0
idna	3.4
importlib-metadata	6.0.0
importlib-resources	5.12.0
iniconfig	2.0.0
ipykernel	6.15.2
ipython	7.34.0
jax	0.3.25

jaxlib	0.3.25
jedi	0.18.2
jupyter_client	7.4.9
jupyter_core	4.11.1
keras	2.11.0
kiwisolver	1.4.4
libclang	15.0.6.1
Markdown	3.4.1
markdown-it-py	2.2.0
MarkupSafe	2.1.2
matplotlib	3.5.3
matplotlib-inline	0.1.6
mccabe	0.7.0
mdurl	0.1.2
mediapy	1.1.2
msgpack	1.0.4
nest-asyncio	1.5.6
numpy	1.21.6
oauthlib	3.2.2
opencv-contrib-python	4.7.0.72
opencv-python	4.7.0.72
opt-einsum	3.3.0
optax	0.1.4
orbax	0.1.0
packaging	23.0
parso	0.8.3
pickleshare	0.7.5
Pillow	9.4.0
pip	22.3.1
platformdirs	3.1.1
pluggy	1.0.0
prompt-toolkit	3.0.38
protobuf	3.19.6
psutil	5.9.0
pyasn1	0.4.8
pyasn1-modules	0.2.8
pycodestyle	2.10.0
pydocstyle	6.2.3
pyflakes	3.0.1
Pygments	2.14.0
pylint	2.17.0
pyparsing	3.0.9
pytest	7.2.2
python-dateutil	2.8.2

python-lsp-jsonrpc	1.0.0
python-lsp-server	1.7.1
pytoolconfig	1.2.5
pywin32	305.1
PyYAML	6.0
pyzmq	23.2.0
rawpy	0.18.0
requests	2.28.2
requests-oauthlib	1.3.1
rich	13.3.1
rope	1.7.0
rsa	4.9
scipy	1.7.3
setuptools	65.6.3
six	1.16.0
snowballstemmer	2.2.0
tensorboard	2.11.2
tensorboard-data-server	0.6.1
tensorboard-plugin-wit	1.8.1
tensorflow	2.11.0
tensorflow-estimator	2.11.0
tensorflow-intel	2.11.0
tensorflow-io-gcs-filesystem	0.31.0
tensorstore	0.1.28
termcolor	2.2.0
tomli	2.0.1
tomlkit	0.11.6
toolz	0.12.0
tornado	6.2
traitlets	5.9.0
typed-ast	1.5.4
typing_extensions	4.5.0
ujson	5.7.0
urllib3	1.26.14
wcwidth	0.2.6
Werkzeug	2.2.3
whatthepatch	1.0.4
wheel	0.38.4
wincertstore	0.2
wrapt	1.15.0
yapf	0.32.0
zipp	3.15.0